

COMP 4601A

Winter 2026 – Assignment #1

Objectives

The goal of this assignment is to implement a basic search engine using the concepts and tools covered in the first half of the course. To complete this assignment, you will need to implement a web crawler, a RESTful search server, and a browser-based client that will allow a user to perform searches. You will also need to deploy your server to OpenStack and integrate it with a distributed search service.

Submission Requirements

There will be several requirements for submission:

1. The code for your assignment must be submitted on Brightspace. You should not submit your data or database files. This submission must include a README file with your name, any partner names (if applicable), a summary of the parts of the assignment that you did/didn't complete successfully, a link to your video demonstration, and the URLs the TA should use to query your search engine (fruits and personal datasets). **If submitting as a group, only one member should make a submission.**
2. Your assignment must be deployed to OpenStack and running throughout the grading process so the TA can execute search requests. Many requests will be made to your server by the TAs, the grading server, and through the distributed search functionality. Some of these may have invalid query parameters or information. **Your server should be implemented in a way that handles any problems gracefully and continues running.** If your server is not running or fails to produce correct results when the TA goes to check your search results, your grade will be significantly affected.
3. Your assignment must be registered with the A1 test through the lab grader and be running throughout the grading process. Search queries will be sent to your server from the grader. To pass tests, your server must respond with proper results within 1 second.
4. You must record a <15-minute video where you:
 - a. Demonstrate your search engine's functionality for both crawled sites.
 - b. Discuss the design of your implementation (see the end of the document for some ideas of things you could discuss).
 - c. Discuss the personal site you selected, the challenges that it presented, and how you addressed these challenges.
5. Your video demonstration must be hosted online. You can use any tools/host you want. Carleton offers support for recording and hosting videos using Kaltura and MediaSpace. See <https://carleton.ca/kaltura/personal/record-with-personal-capture/> and <https://carleton.ca/kaltura/add-new-media/> for information. If you are using MediaSpace, you can set your video to 'Unlisted' and share the link. Ensure the video can be viewed even if you are not logged into your account.

Assignment Requirements

The web crawler portion of your assignment must be capable of crawling the following:

1. The fruit example site. Start at <https://people.scs.carleton.ca/~avamckenney/fruitsA/N-0.html> and crawl the entire site (100 pages).
2. Another site of your choosing. The total number of crawled pages must be at least 500. A recommended maximum page cap is 2500 but you can crawl as many as you desire. It is suggested to limit your crawl within the same domain that you begin. You can design your selection policy to focus on any pages or resources you deem important. You are not required to crawl non-HTML resources but can choose to do so.

Your crawled data must be stored in a database for persistence. You must also perform PageRank calculations and store the values for each page in the database.

Your RESTful web server must read the data from the database, perform required indexing, and provide relevant, ranked search results for any valid request. Your server must support GET requests for the following endpoints:

1. /fruitsA – represents a request to search the data from the fruit example
2. /personal – represents a request to search the data in the alternate site you selected

Both of your search endpoints (/fruits and /personal) must support all combinations of the following query parameters:

1. q – A string representing the search query the user has entered, which may contain multiple words.
2. boost – Either true or false, indicating whether each page should be boosted in the search results using its PageRank score (defaults to false if not provided).
3. limit – A number specifying how many results the user wants returned (minimum 1, maximum 50, default 10). If a valid limit parameter X is specified, your server MUST return X results, even if all documents have a score of 0 (return any X documents in this case).

The browser-based interface for searching must allow the client to specify:

1. The text for their search
2. Whether they want the results to be boosted or not using PageRank
3. The number of results they want to receive (minimum 1, maximum 50, default 10)

The search results displayed in the browser must contain:

1. The URL to the original page
2. The title of the original page
3. The computed search score for the page
4. The PageRank of the page within your crawled network
5. A link to view the data your search engine has for this page. This must include at least the URL, title, list of incoming links to this page, list of outgoing links from this page, and

word frequency information for the page (e.g., banana occurred 6 times, apple occurred 9 times, etc.). You can also display any additional data you produced during the crawl.

If a search request specifies that JSON data should be sent in the response, the response should be a JSON string with the format { "result" : [arrayOfResults] }. The array of results must contain the most relevant *X* results, ranked from highest to lowest score. Each object entry in the results array should contain the following keys:

1. url – The URL of the original page
2. score – The page's computed search score for this search request
3. title - The title of the original page
4. pr - The PageRank of the page within your crawled network

Potential Discussion Points

Below are some things to consider addressing in your demonstration video. This list is not exhaustive, so include any other meaningful analysis you think is valuable.

1. How does your crawler work? What information does it extract from the page? How does it store the data? Is there any intermediary processing you perform to facilitate the later steps of the assignment?
2. Discuss the RESTful design of your server. How has your implementation incorporated the various REST principles?
3. How have you addressed the issue of scalability within the implementation of your crawler and search engine?
4. Explain how the content score for the search is generated.
5. Discuss the PageRank calculation and how you have implemented it.
6. How have you defined your page selection policy for your crawler for your personal site?
7. Why did you select the personal site you chose? Did you run into any problems when working with this site? How did you address these problems?
8. Critique your search engine. How well does it work? How well will it scale? How well will it work with other sites? How do you think it could be improved?